



SLURM: Simple Linux Utility for Resource Management

*ClusterWorld Conference and Expo
June 2003*

Morris Jette and Mark Grondona
Lawrence Livermore National Laboratory

www.llnl.gov/linux/slurm



Disclaimer



This document was prepared as an account of work sponsored by an agency of the United States Government. Neither the United States Government nor the University of California nor any of their employees, makes any warranty, express or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. References herein to any specific commercial product, process, or service by trade name, trademark, manufacture, or otherwise, does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States Government or the University of California. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States Government or the University of California, and shall not be used for advertising or product endorsement purposes.

This work was performed under the auspices of the U.S. Department of Energy by the University of California, Lawrence Livermore National Laboratory under Contract No. W-7405-Eng-48.



Acknowledgments



- > Jointly developed by
 - Lawrence Livermore National Laboratory (LLNL) and
 - Linux NetworX

- > Additional Developers:
 - Jay Windley, Linux NetworX
 - Joseph Ekstrom, LLNL
 - James Garlick , LLNL
 - Kevin Tew , LLNL



What is SLURM?

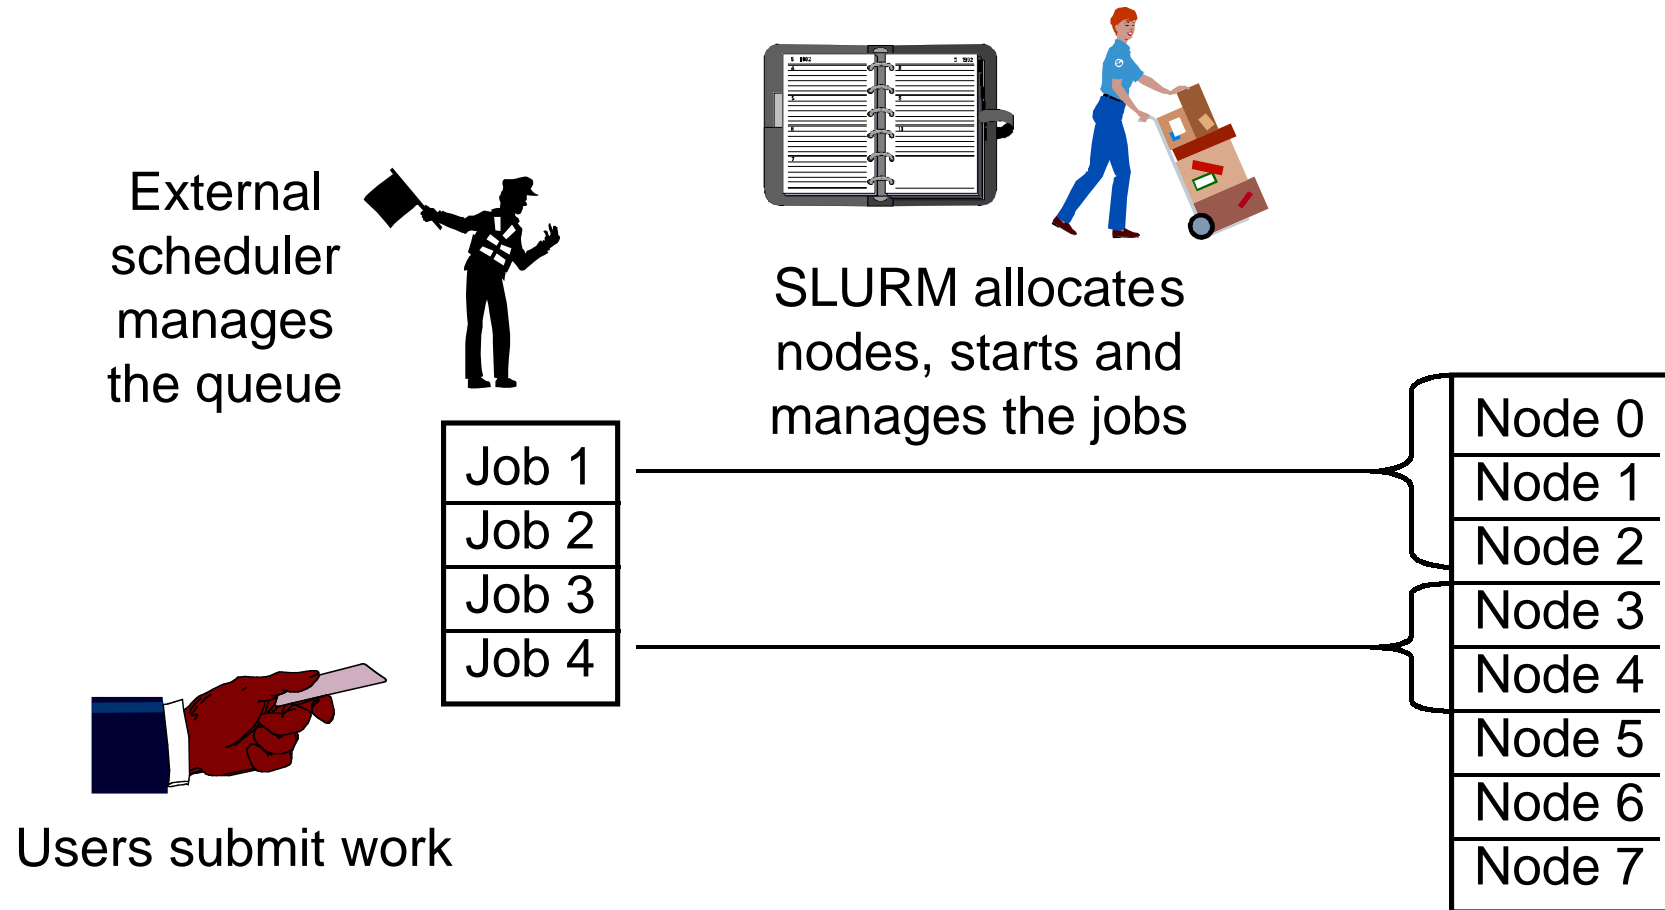


- > Allocates access to computer nodes
- > Distributes work to allocated resources
- > Arbitrates requests by managing queue of pending work

- > NOT a comprehensive cluster administration or monitoring package
- > NOT a sophisticated batch system
 - An external entity can manage the SLURM queues



SLURM in a Nutshell





SLURM Design Criteria



- > Simple
- > Open source: GPL
- > Portable
 - C-language, *autoconf*, general-purpose plugin mechanism
- > Fault-tolerant
 - For SLURM daemons and (optionally) its jobs
- > Secure
- > System administrator friendly
 - Simple configuration file, supports heterogeneous clusters
- > Scalable to thousands of nodes



SLURM Plugins



- > Dynamically linked objects loaded at run time per configuration file
- > Authentication
 - Authd, Munge, or none
- > Interconnect
 - UDP/IP, Quadrics Elan3, or Myrinet
- > Scheduler
 - Maui or FIFO

SLURM daemons and commands

Authentication

Interconnect

Scheduler

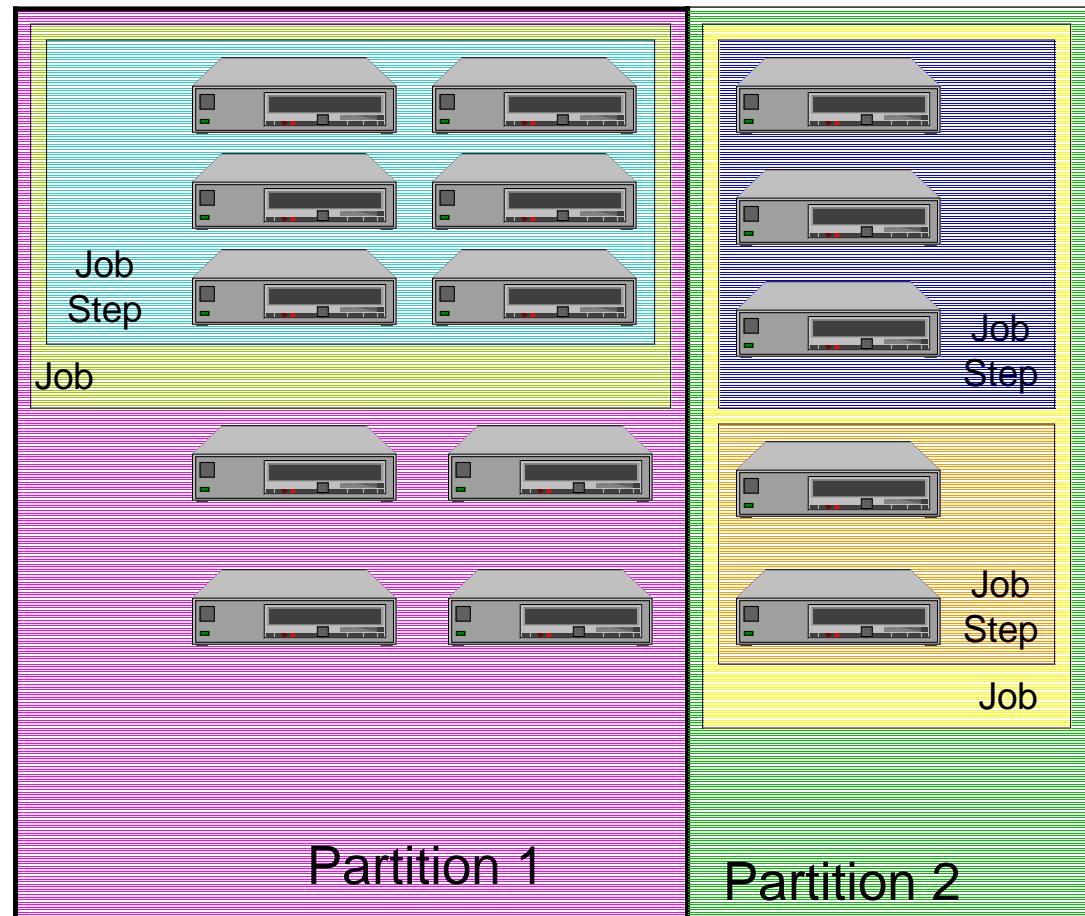
Others...



SLURM Entities



- > Nodes
- > Partitions
- > Jobs
- > Job steps





SLURM Architecture



- > Two daemons
 - slurmctld - controller, optional backup
 - slurmd - computer node daemon

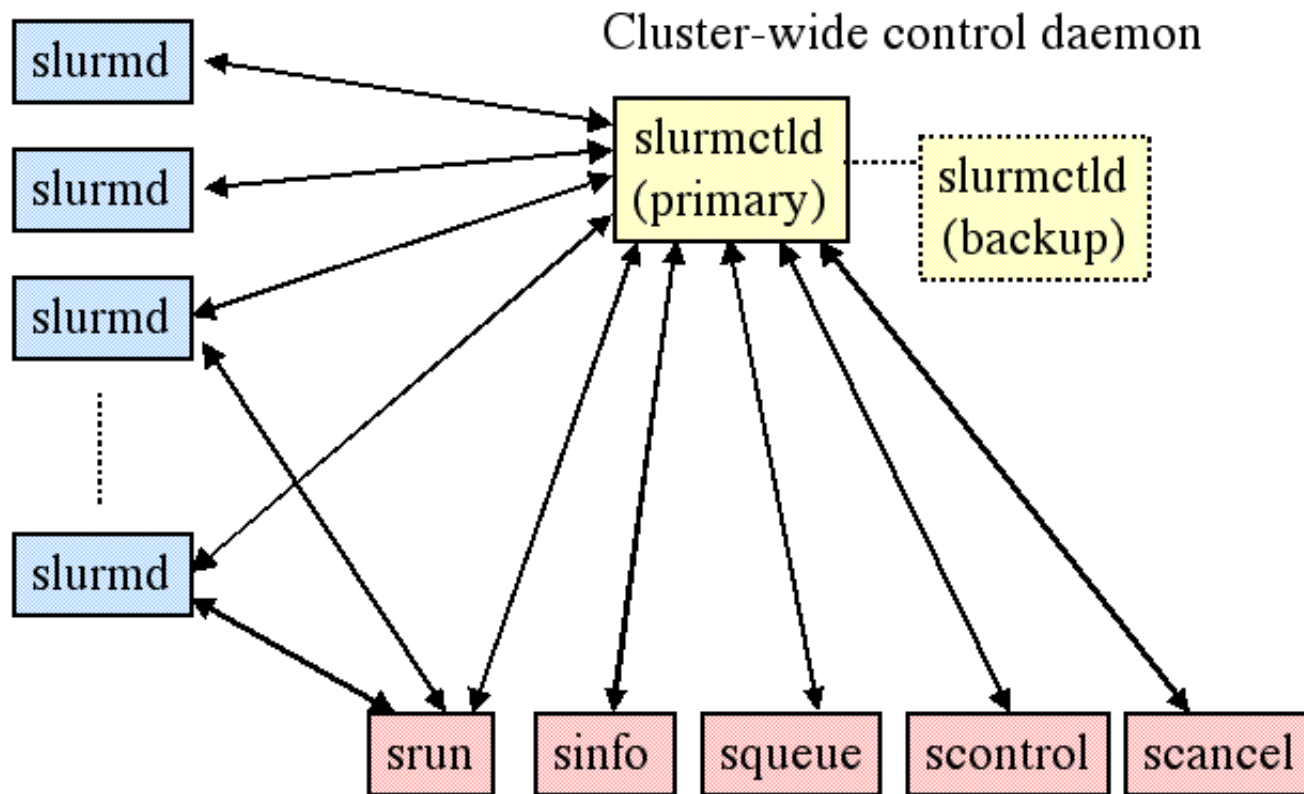
- > Five user commands
 - scontrol - administration tool, get/set configuration
 - sinfo - reports general system information
 - squeue - reports job and job step information
 - srun - submit/initiate job or job step
 - scancel - signal or cancel a job or job step



SLURM Architecture



One daemon per node



User and administrator tools



slurmctld



- > Controller of SLURM (with optional backup)
- > Multi-threaded
- > Independent read and write locks by data structure
- > Nodes represented with bit-maps for rapid manipulations

- > Components
 - Node Manager - node state information
 - Partition Manager - allocates nodes
 - Job Manager - manages queue of pending jobs



slurmd



- > Daemon executing on each compute node
- > Minimal state information
- > Performs actions as directed by slurmctld and srun

- > Components
 - Machine Status
 - Job Status
 - Remote Execution
 - Stream Copy (stdin, stdout, and stderr)
 - Job Control



srun

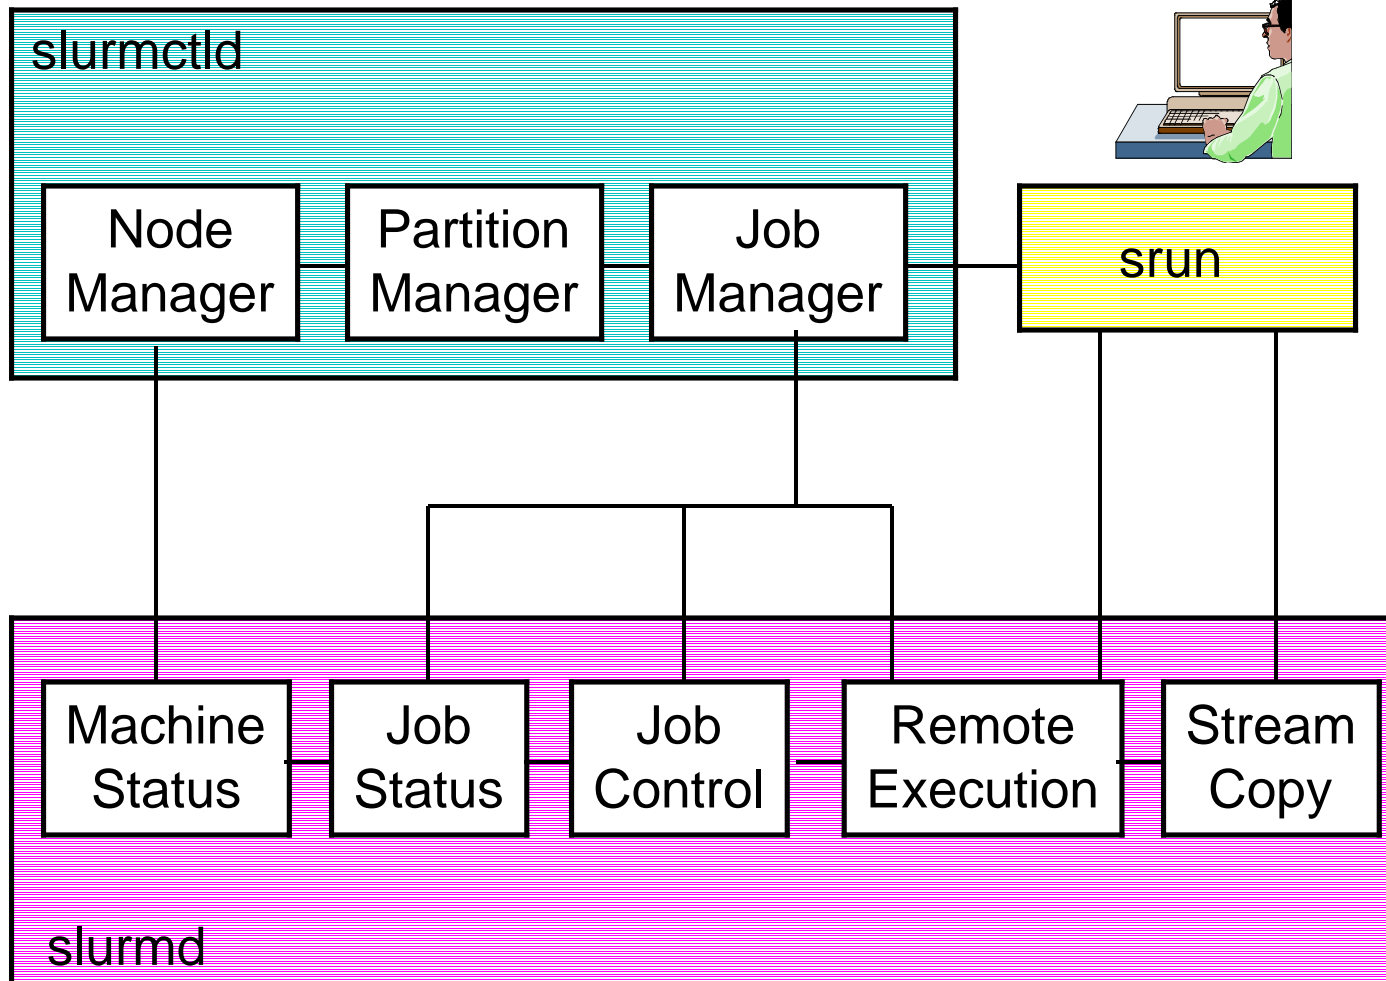


- > User tool to initiate jobs and job steps
 - Allocate resources
 - Submit batch jobs
 - Run jobs interactively
 - Attach to currently running job
 - Launch a set of parallel tests (job step)

- > 13 options to specify resource requirements
 - Partition, processor count, node count, minimum memory per node, minimum processor count per node, specific nodes to use or avoid, node can be share, etc.




SLURM Subsystems





Interactive Job Initiation Overview



- 1: `srun` connects to `slurmctld` requesting resource allocation and step creation
- 2: `slurmctld` responds with node list and job step credential
- 3: `srun` opens I/O connection (ephemeral port)
- 4: `srun` sends job step requests to `slurmd` daemons
- 5: `slurmd` initiates job step and makes I/O connections to `srun`



Sample SLURM Configuration

(excerpt)




```
# Sample SLURM configuration (excerpt)
ControlMachine=linux0
BackupController=linux1
#
AuthType="auth/authd"
HeartbeatInterval=60
PluginDir=/usr/lib/slurm
SlurmctlPort=7002
SlurmdPort=7003
SlurmUser=slurm
#
NodeName=DEFAULT Procs=2 TmpDisk=64000
NodeName=linux[2-1000]    RealMemory=16000 Weight=16
NodeName=linux[1001-1016] RealMemory=32000 Weight=32
#
PartitionName=debug Nodes=linux[2-33]    MaxTime=30
PartitonName=batch  Nodes=linux[34-1016] MaxTime=Infinite
```




Performance Results

/bin/hostname, 2 tasks per node





SLURM Plans



- > Support more systems
 - Infiniband, IBM Blue Gene/L
- > Integrate slurmctld with relational database
- > Convert daemon communications to plugin
 - Support broadcast (e.g. STORM)
- > Job preempt/resume
- > Job checkpoint/restart